

Import pandas and numpy

```
In [1]: import numpy as np
import pandas as pd
```

```
In [2]: kcd=pd.read_csv("/home/placement/Downloads/fiat500.csv")
kcd.info
```

```
Out[2]: <bound method DataFrame.info of
wners \
0      1  lounge      51      882  25000      1
1      2   pop      51     1186  32500      1
2      3  sport      74     4658 142228      1
3      4  lounge      51     2739 160000      1
4      5   pop      73     3074 106880      1
...    ...    ...    ...    ...    ...
1533 1534  sport      51     3712 115280      1
1534 1535  lounge      74     3835 112000      1
1535 1536   pop      51     2223  60457      1
1536 1537  lounge      51     2557  80750      1
1537 1538   pop      51     1766  54276      1

      lat      lon  price
0  44.907242  8.611560  8900
1  45.666359 12.241890  8800
2  45.503300 11.417840  4200
3  40.633171 17.634609  6000
4  41.903221 12.495650  5700
...    ...    ...    ...
1533 45.069679  7.704920  5200
1534 45.845692  8.666870  4600
1535 45.481541  9.413480  7500
1536 45.000702  7.682270  5990
1537 40.323410 17.568270  7900

[1538 rows x 9 columns]>
```

In [3]: kcd

Out[3]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price
0	1	lounge	51	882	25000	1	44.907242	8.611560	8900
1	2	pop	51	1186	32500	1	45.666359	12.241890	8800
2	3	sport	74	4658	142228	1	45.503300	11.417840	4200
3	4	lounge	51	2739	160000	1	40.633171	17.634609	6000
4	5	pop	73	3074	106880	1	41.903221	12.495650	5700
...
1533	1534	sport	51	3712	115280	1	45.069679	7.704920	5200
1534	1535	lounge	74	3835	112000	1	45.845692	8.666870	4600
1535	1536	pop	51	2223	60457	1	45.481541	9.413480	7500
1536	1537	lounge	51	2557	80750	1	45.000702	7.682270	5990
1537	1538	pop	51	1766	54276	1	40.323410	17.568270	7900

1538 rows × 9 columns

In [4]: a=kcd.groupby(['model']).count()
a

Out[4]:

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	price
model								
lounge	1094	1094	1094	1094	1094	1094	1094	1094
pop	358	358	358	358	358	358	358	358
sport	86	86	86	86	86	86	86	86

```
In [5]: drop=kcd.drop(['ID','lat','lon'],axis=1)
drop
```

Out[5]:

	model	engine_power	age_in_days	km	previous_owners	price
0	lounge	51	882	25000	1	8900
1	pop	51	1186	32500	1	8800
2	sport	74	4658	142228	1	4200
3	lounge	51	2739	160000	1	6000
4	pop	73	3074	106880	1	5700
...
1533	sport	51	3712	115280	1	5200
1534	lounge	74	3835	112000	1	4600
1535	pop	51	2223	60457	1	7500
1536	lounge	51	2557	80750	1	5990
1537	pop	51	1766	54276	1	7900

1538 rows × 6 columns

```
In [6]: drop['model']=drop['model'].map({'lounge':1,'pop':2,'sport':3})
drop
```

Out[6]:


	model	engine_power	age_in_days	km	previous_owners	price
0	1	51	882	25000	1	8900
1	2	51	1186	32500	1	8800
2	3	74	4658	142228	1	4200
3	1	51	2739	160000	1	6000
4	2	73	3074	106880	1	5700
...
1533	3	51	3712	115280	1	5200
1534	1	74	3835	112000	1	4600
1535	2	51	2223	60457	1	7500
1536	1	51	2557	80750	1	5990
1537	2	51	1766	54276	1	7900

1538 rows × 6 columns

```
In [7]: cor1=drop.corr()
```

```
In [8]: import seaborn as sb  
sb.heatmap(corr1,vmax=1,vmin=-1,annot=True,linewidths=5,cmap='bwr')
```

Out[8]: <Axes: >



```
In [9]: y=drop['price']
x=drop.drop('price',axis=1)
x
```

Out[9]:

	model	engine_power	age_in_days	km	previous_owners
0	1	51	882	25000	1
1	2	51	1186	32500	1
2	3	74	4658	142228	1
3	1	51	2739	160000	1
4	2	73	3074	106880	1
...
1533	3	51	3712	115280	1
1534	1	74	3835	112000	1
1535	2	51	2223	60457	1
1536	1	51	2557	80750	1
1537	2	51	1766	54276	1

1538 rows × 5 columns

```
In [10]: y
```

```
Out[10]: 0      8900
1      8800
2      4200
3      6000
4      5700
```

```
...
1533    5200
1534    4600
1535    7500
1536    5990
1537    7900
```

Name: price, Length: 1538, dtype: int64

```
In [11]: !pip3 install scikit-learn
```

```
Requirement already satisfied: scikit-learn in ./local/lib/python3.8/site-packages (1.2.2)  
Requirement already satisfied: scipy>=1.3.2 in ./local/lib/python3.8/site-packages (from scikit-learn) (1.10.1)  
Requirement already satisfied: threadpoolctl>=2.0.0 in ./local/lib/python3.8/site-packages (from scikit-learn) (3.1.0)  
Requirement already satisfied: numpy>=1.17.3 in ./local/lib/python3.8/site-packages (from scikit-learn) (1.24.3)  
Requirement already satisfied: joblib>=1.1.1 in ./local/lib/python3.8/site-packages (from scikit-learn) (1.2.0)
```

```
In [12]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.33,random_state=42)
```

```
In [13]: x_test.head(10)
```

```
Out[13]:
```

	model	engine_power	age_in_days	km	previous_owners
481	2	51	3197	120000	2
76	2	62	2101	103000	1
1502	1	51	670	32473	1
669	1	51	913	29000	1
1409	1	51	762	18800	1
1414	1	51	762	39751	1
1089	1	51	882	33160	1
1507	1	51	701	17324	1
970	1	51	701	29000	1
1198	1	51	1155	38000	1

LinearRegression

```
In [14]: from sklearn.linear_model import LinearRegression
reg=LinearRegression()
reg.fit(x_train,y_train)
```

```
Out[14]: ▼ LinearRegression
LinearRegression()
```

```
In [15]: y_pred=reg.predict(x_test)
y_pred
```

```
8012.59977099, 9749.71169477, 5970.69892919, 10374.19319599,
5505.37212671, 9603.80845104, 10080.60444002, 10173.49549365,
9553.48664879, 4886.66020447, 5826.86758437, 7127.78307449,
9986.09840714, 10375.1819333, 9936.20036211, 7755.07402414,
8820.32822581, 10009.77294687, 10261.12807264, 9955.43106434,
8385.1116117, 9441.36137497, 8621.10384346, 9719.70050582,
9767.12327701, 9755.03033027, 6859.84033207, 7339.68592914,
8740.34003982, 9898.84623968, 9788.7072129, 10439.74281794,
8145.90808395, 6767.15633519, 9962.57850061, 8846.92420399,
9927.58506055, 10279.88133318, 10205.11210182, 10065.46678709,
9343.97683092, 9983.85933876, 9237.93178546, 10073.45985579,
7906.63849672, 6017.75726035, 8780.77873324, 10211.55465771,
5737.35007744, 10190.21750673, 9661.444679, 7747.41088806,
9396.65945773, 7357.03908605, 10261.68730153, 10041.70922157,
10525.09542651, 9941.6915233, 10042.87112799, 6342.10368715,
10588.92756092, 9940.98736563, 10501.95046891, 9697.00608104,
9642.20441674, 6177.49903451, 8056.81304643, 10318.99744586,
6334.90676093, 7347.76781534, 10049.18638926, 6780.85650138,
7897.31981053, 5062.64376289, 4656.55980585, 8690.25433913,
6988.39956167, 7416.44791638, 6784.57575877, 7034.60046808,
```


Efficiency

```
In [16]: from sklearn.metrics import r2_score  
r2_score(y_test,y_pred)
```

```
Out[16]: 0.8383895235218546
```

Mean squared error

```
In [17]: from sklearn.metrics import mean_squared_error as mc  
sq=mc(y_test,y_pred)  
sq
```

```
Out[17]: 593504.2888137395
```

```
In [18]: import math as m  
dp=m.sqrt(sq)  
print(dp)
```

```
770.3922954013361
```

In [19]: y_pred

```
Out[19]: array([ 5994.51703157, 7263.58726658, 9841.90754881, 9699.31627673,
 10014.19892635, 9630.58715835, 9649.4499026 , 10092.9819664 ,
 9879.19498711, 9329.19347948, 10407.2964056 , 7716.91706011,
 7682.89152522, 6673.95810983, 9639.42618839, 10346.53679153,
 9366.53363673, 7707.90063494, 4727.33552438, 10428.17092937,
 10359.87663878, 10364.84674179, 7680.16157493, 9927.58506055,
 7127.7284177 , 9097.51161986, 4929.31229715, 6940.60225317,
 7794.35120591, 9600.43942019, 7319.85877519, 5224.05298205,
 5559.52039134, 5201.35403287, 8960.11762682, 5659.72968338,
 9915.79926869, 8255.93615893, 6270.40332834, 8556.73835062,
 9749.72882426, 6873.76758364, 8951.72659758, 10301.95669828,
 8674.89268564, 10301.93257222, 9165.73586068, 8846.92420399,
 7044.68964545, 9052.4031418 , 9390.75738772, 10267.3912561 ,
 10046.90924744, 6855.71260655, 9761.93338967, 9450.05744337,
 9274.98388541, 10416.00474283, 9771.10646661, 7302.96566423,
 10082.61483093, 6996.96553454, 9829.40534825, 7134.21944391,
 6407.26222178, 9971.82132188, 9757.01618446, 8614.84049875,
 8437.92452169, 6489.24658616, 7752.65456507, 6626.60510856,
 8329.88998217, 10412.00324329, 7342.77348105, 8543.63624413,
 8706.44742777, 10010.42502651, 7256.06706062, 8522.1400051 ])
```

```
In [20]: results=pd.DataFrame(columns=['price','predicted'])
results['price']=y_test
results['predicted']=y_pred
results.head(10)
```

Out[20]:

	price	predicted
481	7900	5994.517032
76	7900	7263.587267
1502	9400	9841.907549
669	8500	9699.316277
1409	9700	10014.198926
1414	9900	9630.587158
1089	9900	9649.449903
1507	9950	10092.981966
970	10700	9879.194987
1198	8999	9329.193479

```
In [21]: results['actual price']=results.apply(lambda column:column.price-column.predicted,axis=1)  
results
```

Out[21]:

	price	predicted	actual price
481	7900	5994.517032	1905.482968
76	7900	7263.587267	636.412733
1502	9400	9841.907549	-441.907549
669	8500	9699.316277	-1199.316277
1409	9700	10014.198926	-314.198926
...
291	10900	10007.364639	892.635361
596	5699	6390.174715	-691.174715
1489	9500	10079.478928	-579.478928
1436	6990	8363.337585	-1373.337585
575	10900	10344.486077	555.513923

508 rows × 3 columns

In []: